

# RATS

## Rats: hierarchical linear regression

This example is taken from the openBUGS example library, originally from section 6 of Gelfand et al (1990), and concerns 30 young rats whose weights were measured weekly for five weeks.

The goal of this example is to give everyone a look at what fitting a model in JAGS looks like, (the nuts and bolts) but also a little bit of hierarchical modelling. We'll start with a single regression model, and then add one level of hierarchy: a population intercept and a population slope.

Let's look at the data, which we first have to load:


```
Weights <- read.csv("GiantRats.csv",header=T)
head(Weights)
```

```
##      week1 week2 week3 week4 week5
## 1      151    199    246    283    320
## 2      145    199    249    293    354
## 3      147    214    263    312    328
## 4      155    200    237    272    297
## 5      135    188    230    280    323
## 6      159    210    252    298    331
```


```
data <- list(x = c(8.0, 15.0, 22.0, 29.0, 36.0), xbar = 22, N = 30, T = 5,
            Y = Weights)
```

Here, x is time (days), N is the number of rats, T is the number of time points (5), and Y are the rat weights.

This is a single rat's weight over the 5 week period.



You can see there is some variation for the other rats, but a similar shape



## OK, so let's look at the model

We would like to quantify the mean rate of growth of the rats per week, that is, we want to quantify the slope. So, we will fit a simple linear regression.

First is the likelihood and mean model:

**none of this model code is runnable here! !**

It's located in a separate file "RATS\_model.R"

```

model{ #all of this is wrapped in a "model" loop
  # specify likelihood of each data component by looping through all
  # elements of the data (weights, Y), which is read-in as a matrix,
  # so need to loop through each row (rat) and each column (timepoints):
  for(i in 1:N){ #loop across rats
    for(j in 1:T){ #loop across time points (weeks)
      Y[i,j] ~ dnorm(mu[i,j], tau) # this is the likelihood (distribution of observations)
      #we choose a normal likelihood bc assume data are normally distributed.
      mu[i,j] <- alpha + beta*(x[j] - xbar) # this is the regression ("mean model")
    }
  }
  ....
}

```

Next, we have to assign prior distributions to the stochastic parameters

```

....
# Conjugate, relatively non-informative priors to root
# nodes (population parameters)
alpha ~ dnorm(0,0.00001)
beta ~ XXX #what goes here?
tau ~ dgamma(0.01,0.01) #this is a precision (simply the inverse of the variance)

sigma <- 1/sqrt(tau) # convert the precision to a standard deviation
}

```

**beta ~ ??? --> Please take a minute to think about what this should be.**

**What is beta?**

**let's look at the whole model**

```

model{
  #likelihood and mean model
  for(i in 1:N){
    for(j in 1:T){
      Y[i,j] ~ dnorm(mu[i,j], tau)
      mu[i,j] <- alpha + beta*(x[j] - xbar)
    }
  }
  #prior distributions
  alpha ~ dnorm(0,0.00001)
  beta ~ dnorm(0,0.00001)
  tau ~ dgamma(0.01, 0.01) #this is called a "confugate prior". Allows gibbs sampling
  .
  sigma <- 1/sqrt(tau)
}

```

# Before we run the model, let's add a level of hierarchy

Want to quantify growth rates and intercepts for each rat, as well as overall means. Random-intercepts, random-slopes model.

```

model{
  for(i in 1:N){
    for(j in 1:T){
      Y[i,j] ~ dnorm(mu[i,j], tau)
      mu[i,j] <- alpha[i] + beta[i]*(x[j] - xbar) #notice, we now have one intercept
and slope for each rat, i.
    }
  }
  ...
}

```

So what do priors look like now?

```

...
#prior distributions#
  #hierarchical priors
  for(i in 1:N){
    alpha[i] ~ dnorm(mu.alpha, tau.alpha)
    beta[i] ~ ???
  }
  #root, or global priors
mu.alpha ~ dnorm(0,0.001)
mu.beta ~ ???
  #precisions
tau ~ dgamma(0.01,0.01)
tau.alpha ~ ???
tau.beta ~ ???

# Calculate standard deviations associated with each
# precition:
sigma <- 1/sqrt(tau)
sigma.alpha <- ???
sigma.beta <- ???

}

```

# Everyone take a minute or two to think about how to fill in the blanks (???)

## Any questions so far?

This model (above) will be saved in a .R file called "RATS\_Hmodel.R". There is also a "RATS\_model.R" which is non-hierarchical, the first version we looked at. We can call either model in as part of the JAGS model run.

## Final step before we run the model

Need to provide intial values to the stochastic parameters (anything with a prior distribution). #Which parameters need initial values? Specify initials for root nodes, for 3 MCMC chains as a list of lists:

```

inits = list(
  list(mu.alpha = 1, mu.beta = 0, tau=1,tau.alpha=1,tau.beta=1),
  list(mu.alpha = xx, mu.beta = xx, tau=0.5,tau.alpha=2,tau.beta=0.5),
  list(mu.alpha = xx, mu.beta = xx, tau=2,tau.alpha=0.5,tau.beta=2))

```

```
inits = list(  
  list(mu.alpha = 1, mu.beta = 0, tau=1,tau.alpha=1,tau.beta=1),  
  list(mu.alpha = 5, mu.beta = 0, tau=0.5,tau.alpha=2,tau.beta=0.5),  
  list(mu.alpha = 0.5, mu.beta = 0, tau=2,tau.alpha=0.5,tau.beta=2))
```

# Run the hierarchical model

Need to provide the function with data, initial values, number of chains (more on this shortly), and adaptive iterations. #Let's look at how that works first The model will run three parallel MCMC chains. MCMC stands for Markov chain Monte Carlo, ie, the chains proceed at each iteration by random sampling from distributions, followed by acceptance or rejection of proposed samples, and adjustment of the sampling distributions. This is akin to maximum likelihood estimation, except that the sampling algorithms "under the hood" vary according to the model selected. Different programs openBUGS, JAGS, Stan, have different available sampling algorithms. These are, for the most part, not something the user interacts with in any way, that is, this is automated.

Here, we will just define the model and sample it for 10 iterations (Note the n.adapt here is way too short, but this is useful for this example)

```
library(rjags) #this package interfaces between R and JAGS
```

```
## Loading required package: coda
```

```
## Linked to JAGS 4.2.0
```


```
## Loaded modules: basemod,bugs
```

```
Rats_model <- jags.model("RATS_Hmodel.R", data=data, inits=inits, n.chains = 3, n.adapt=1)
```


```
## Compiling model graph  
##   Resolving undeclared variables  
##   Allocating nodes  
## Graph information:  
##   Observed stochastic nodes: 150  
##   Unobserved stochastic nodes: 65  
##   Total graph size: 550  
##  
## Initializing model
```

```
burnin_samples <- coda.samples(Rats_model,variable.names = c("mu.beta"),n.iter=10)  
plot(burnin_samples)
```

### Trace of mu.beta



### Density of mu.beta



#This model has not converged, yet. What you see on the left is the value of the " $\mu.\text{beta}$ " parameter for each of three chains (black, red, green) at each iteration up to 10 along the x-axis. This is a trace plot, or history plot, and shows how the chains are moving through parameter space.

On the right is the density plot, a histogram which shows the distribution of samples for  $\mu.\text{beta}$ .

Two things show the model is not yet converged: 1) Depending on which iterations along the x-axis you look at, the mean of all three is different. This is shown by a wobbly density plot. 2) Different chains are doing different things. They don't agree. This is why we start the three chains in different places.

## Let's run the model longer:

(Note the much more realistic  $n.\text{adapt}=1000$ )


```
Rats_model <- jags.model("RATS_Hmodel.R", data=data, inits=inits, n.chains = 3, n.adapt=500)
```




```
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 150
##   Unobserved stochastic nodes: 65
##   Total graph size: 550
##
## Initializing model
```

```
converged_samples <- coda.samples(Rats_model,variable.names = c("mu.beta"),n.iter=1000)
plot(converged_samples) #base R method to plot output. Flexible but less user-friendly.
```

**Trace of mu.beta**



**Density of mu.beta**



#this model is now converged. We ran many more samples, but you'll notice that the mean is stable across iterations, and the chains agree on the mean (and uncertainty).

# Let's sample the rest of the parameters, and let's monitor the deviance.

```
load.module("dic") #loads the DIC module
```

```
## module dic loaded
```

```
full_coda <- coda.samples(Rats_model, variable.names = c("deviance", "mu.alpha", "mu.beta", "alpha", "beta", "sig", "sig.alpha", "sig.beta"), n.iter=1000)  
library(mcmcplots) #another way to plot output, more user-friendly, but inflexible. A nnoying for large models.  
mcmcplot(full_coda)
```

```
##  
  
Preparing plots for deviance. 2% complete.
```

```
##  
  
Preparing plots for mu.alpha. 3% complete.
```

```
##  
  
Preparing plots for mu.beta. 5% complete.
```

```
##  
  
Preparing plots for sig. 6% complete.
```

```
##  
  
Preparing plots for sig.alpha. 8% complete.
```

```
##  
  
Preparing plots for sig.beta. 9% complete.
```

```
##  
  
Preparing plots for alpha. 11% complete.
```

```
##  
  
Preparing plots for alpha. 12% complete.
```

##

Preparing plots for alpha. 14% complete.

##

Preparing plots for alpha. 15% complete.

##

Preparing plots for alpha. 17% complete.

##

Preparing plots for alpha. 18% complete.

##

Preparing plots for alpha. 20% complete.

##

Preparing plots for alpha. 21% complete.

##

Preparing plots for alpha. 23% complete.

##

Preparing plots for alpha. 24% complete.

##

Preparing plots for alpha. 26% complete.

##

Preparing plots for alpha. 27% complete.

##

Preparing plots for alpha. 29% complete.

##

Preparing plots for alpha. 30% complete.

##

Preparing plots for alpha. 32% complete.

##

Preparing plots for alpha. 33% complete.

##

Preparing plots for alpha. 35% complete.

##

Preparing plots for alpha. 36% complete.

##

Preparing plots for alpha. 38% complete.

##

Preparing plots for alpha. 39% complete.

##

Preparing plots for alpha. 41% complete.

##

Preparing plots for alpha. 42% complete.

##

Preparing plots for alpha. 44% complete.

##

Preparing plots for alpha. 45% complete.

##

Preparing plots for alpha. 47% complete.

##

Preparing plots for alpha. 48% complete.

##

Preparing plots for alpha. 50% complete.

##

Preparing plots for alpha. 52% complete.

##

Preparing plots for alpha. 53% complete.

##

Preparing plots for alpha. 55% complete.

##

Preparing plots for beta. 56% complete.

##

Preparing plots for beta. 58% complete.

##

Preparing plots for beta. 59% complete.

##

Preparing plots for beta. 61% complete.

##

Preparing plots for beta. 62% complete.

##

Preparing plots for beta. 64% complete.

##

Preparing plots for beta. 65% complete.

##

Preparing plots for beta. 67% complete.

##

Preparing plots for beta. 68% complete.

##

Preparing plots for beta. 70% complete.

##

Preparing plots for beta. 71% complete.

##

Preparing plots for beta. 73% complete.

##

Preparing plots for beta. 74% complete.

##

Preparing plots for beta. 76% complete.

##

Preparing plots for beta. 77% complete.

##

Preparing plots for beta. 79% complete.

##

Preparing plots for beta. 80% complete.

##

Preparing plots for beta. 82% complete.

##

Preparing plots for beta. 83% complete.

##

Preparing plots for beta. 85% complete.

##

Preparing plots for beta. 86% complete.

##

Preparing plots for beta. 88% complete.

```
##
```

```
Preparing plots for beta. 89% complete.
```

```
##
```

```
Preparing plots for beta. 91% complete.
```

```
##
```

```
Preparing plots for beta. 92% complete.
```

```
##
```

```
Preparing plots for beta. 94% complete.
```

```
##
```

```
Preparing plots for beta. 95% complete.
```

```
##
```

```
Preparing plots for beta. 97% complete.
```

```
##
```

```
Preparing plots for beta. 98% complete.
```

```
##
```

```
Preparing plots for beta. 100% complete.
```

# This is the bgr (brooks-gelman-rubin) convergence diagnostic

Want to be less than 1.2. But, doesn't work so well with large, complex models. Need to visually evaluate!!

```
gelman.diag(full_coda)
```

```
## Potential scale reduction factors:
```

```
##
```



##	Point	est.	Upper	C.I.
##	alpha[1]	1.00	1.00	
##	alpha[2]	1.00	1.00	
##	alpha[3]	1.00	1.00	
##	alpha[4]	1.00	1.00	
##	alpha[5]	1.00	1.00	
##	alpha[6]	1.00	1.00	
##	alpha[7]	1.00	1.00	
##	alpha[8]	1.00	1.00	
##	alpha[9]	1.00	1.00	
##	alpha[10]	1.00	1.00	
##	alpha[11]	1.00	1.00	
##	alpha[12]	1.00	1.00	
##	alpha[13]	1.00	1.01	
##	alpha[14]	1.00	1.00	
##	alpha[15]	1.00	1.00	
##	alpha[16]	1.00	1.00	
##	alpha[17]	1.00	1.00	
##	alpha[18]	1.00	1.00	
##	alpha[19]	1.00	1.00	
##	alpha[20]	1.00	1.00	
##	alpha[21]	1.00	1.00	
##	alpha[22]	1.00	1.00	
##	alpha[23]	1.00	1.00	
##	alpha[24]	1.00	1.00	
##	alpha[25]	1.00	1.00	
##	alpha[26]	1.00	1.00	
##	alpha[27]	1.00	1.00	
##	alpha[28]	1.00	1.00	
##	alpha[29]	1.00	1.00	
##	alpha[30]	1.00	1.00	
##	beta[1]	1.00	1.00	
##	beta[2]	1.00	1.00	
##	beta[3]	1.00	1.00	
##	beta[4]	1.00	1.00	
##	beta[5]	1.00	1.00	
##	beta[6]	1.00	1.00	
##	beta[7]	1.00	1.00	
##	beta[8]	1.00	1.00	
##	beta[9]	1.00	1.01	
##	beta[10]	1.00	1.00	
##	beta[11]	1.00	1.00	
##	beta[12]	1.00	1.00	
##	beta[13]	1.00	1.00	
##	beta[14]	1.00	1.00	
##	beta[15]	1.00	1.00	
##	beta[16]	1.00	1.01	
##	beta[17]	1.00	1.00	
##	beta[18]	1.00	1.00	
##	beta[19]	1.00	1.00	

```
## beta[20]      1.00      1.01
## beta[21]      1.00      1.00
## beta[22]      1.00      1.00
## beta[23]      1.00      1.01
## beta[24]      1.00      1.00
## beta[25]      1.00      1.01
## beta[26]      1.00      1.00
## beta[27]      1.00      1.01
## beta[28]      1.00      1.00
## beta[29]      1.00      1.01
## beta[30]      1.00      1.00
## deviance      1.00      1.00
## mu.alpha      1.00      1.00
## mu.beta       1.00      1.00
## sig           1.00      1.00
## sig.alpha     1.00      1.00
## sig.beta      1.01      1.02
##
## Multivariate psrf
##
## 1.02
```

## Summarize our model output and get posterior stats (means, sds, 95% CI)

Note that this only summarizes parameters that you monitored (with the `coda.samples()` function above). If you forgot to monitor something.... need to start over with `jags.model()` step.

```
table1=summary(full_coda)$stat
table2=summary(full_coda)$quantiles
outstats<-cbind(table2[,1],table1[,2],table2[,1],table2[,5])
colnames(outstats)<-c("mean","sd","val2.5pc","val97.5pc")
outstats
```

```
##           mean           sd    val2.5pc    val97.5pc
## alpha[1] 234.7292329 2.66122741 234.7292329 245.1797718
## alpha[2] 242.2881352 2.73808530 242.2881352 253.0610466
## alpha[3] 247.3355006 2.73202207 247.3355006 257.7173140
## alpha[4] 227.1029854 2.72928920 227.1029854 237.8190855
## alpha[5] 225.9671306 2.75381648 225.9671306 236.8998613
## alpha[6] 244.6775467 2.62401372 244.6775467 254.9156265
## alpha[7] 223.2005220 2.73326886 223.2005220 233.9118518
## alpha[8] 243.0112633 2.66671216 243.0112633 253.5391654
## alpha[9] 277.7873140 2.69501151 277.7873140 288.3867119
## alpha[10] 213.9474722 2.71356707 213.9474722 224.4991750
## alpha[11] 253.0006709 2.62373509 253.0006709 263.3342949
## alpha[12] 222.7562996 2.71914736 222.7562996 233.5711901
```

##	alpha[13]	237.2714520	2.71416731	237.2714520	247.7262949
##	alpha[14]	262.9393513	2.68366022	262.9393513	273.6234543
##	alpha[15]	237.2273316	2.74755383	237.2273316	248.0224343
##	alpha[16]	239.8754669	2.72494947	239.8754669	250.4040601
##	alpha[17]	226.8752988	2.69265247	226.8752988	237.3834669
##	alpha[18]	234.9620442	2.74031520	234.9620442	245.8591796
##	alpha[19]	248.2654857	2.70608697	248.2654857	259.0939784
##	alpha[20]	236.4082566	2.60484873	236.4082566	247.0577309
##	alpha[21]	243.2003827	2.68902055	243.2003827	253.6477707
##	alpha[22]	219.9092387	2.67410214	219.9092387	230.2988756
##	alpha[23]	223.4184813	2.72384962	223.4184813	233.9614597
##	alpha[24]	239.7935567	2.66864111	239.7935567	250.3331635
##	alpha[25]	228.9133291	2.73062190	228.9133291	239.6820443
##	alpha[26]	248.6417511	2.71566396	248.6417511	259.1234224
##	alpha[27]	249.0128106	2.67800865	249.0128106	259.4249678
##	alpha[28]	237.7650251	2.68883464	237.7650251	248.3198628
##	alpha[29]	212.5082994	2.66142964	212.5082994	223.0781005
##	alpha[30]	236.0840993	2.71145351	236.0840993	246.8340682
##	beta[1]	5.5571876	0.24707128	5.5571876	6.5376317
##	beta[2]	6.5684258	0.25177044	6.5684258	7.5435615
##	beta[3]	5.9993730	0.25078210	5.9993730	6.9825562
##	beta[4]	4.8323989	0.26218558	4.8323989	5.8548588
##	beta[5]	6.0949890	0.25327224	6.0949890	7.0936952
##	beta[6]	5.7215741	0.23682190	5.7215741	6.6396445
##	beta[7]	5.4983458	0.24239638	5.4983458	6.4534387
##	beta[8]	5.9366634	0.24312500	5.9366634	6.8835385
##	beta[9]	6.5581995	0.25561055	6.5581995	7.5443898
##	beta[10]	5.3808786	0.24290007	5.3808786	6.3232553
##	beta[11]	6.3274144	0.25074094	6.3274144	7.2857103
##	beta[12]	5.6365518	0.24097397	5.6365518	6.5895910
##	beta[13]	5.6911477	0.24431803	5.6911477	6.6514707
##	beta[14]	6.2088098	0.24616782	6.2088098	7.1725438
##	beta[15]	4.9291230	0.24525221	4.9291230	5.8897687
##	beta[16]	5.4508246	0.24459439	5.4508246	6.3875205
##	beta[17]	5.8153125	0.23957649	5.8153125	6.7715744
##	beta[18]	5.3644585	0.24808929	5.3644585	6.3213718
##	beta[19]	5.9311952	0.24357986	5.9311952	6.8997346
##	beta[20]	5.5822730	0.23679069	5.5822730	6.5182397
##	beta[21]	5.9163668	0.24330351	5.9163668	6.8729659
##	beta[22]	5.3710181	0.24334345	5.3710181	6.3291709
##	beta[23]	5.2794454	0.24136819	5.2794454	6.2207315
##	beta[24]	5.4027771	0.24368344	5.4027771	6.3695781
##	beta[25]	6.3848664	0.25750166	6.3848664	7.4049201
##	beta[26]	6.0706435	0.24460868	6.0706435	7.0443348
##	beta[27]	5.4240221	0.24446013	5.4240221	6.3812131
##	beta[28]	5.3776202	0.24558422	5.3776202	6.3266535
##	beta[29]	5.1927330	0.24840243	5.1927330	6.1776859
##	beta[30]	5.6441150	0.23932569	5.6441150	6.5888626
##	deviance	941.4167382	14.60427343	941.4167382	997.9925425
##	mu.alpha	234.9071578	2.85235423	234.9071578	246.2754220

## mu.beta	5.9681499	0.11109175	5.9681499	6.4134986
## sig	5.2606387	0.45966951	5.2606387	7.0718542
## sig.alpha	11.3385236	2.03898976	11.3385236	19.2194718
## sig.beta	0.3563018	0.09056212	0.3563018	0.7197392